

# Analysis of Yelp Business Intelligence Data

I will analyze a subset of Yelp's business, reviews and user data. This dataset comes to us from [Kaggle](#) although we have taken steps to pull this data into a public s3 bucket: `s3://sta9760-lei-yelpdatasets/yelp/*.json`

## Installation and Initial Setup

Begin by installing the necessary libraries that you may need to conduct your analysis. At the very least, you must install `pandas` and `matplotlib`

In [1]:

```
%%info
```

```
Current session configs: {'conf': {'spark.pyspark.python': 'python3', 'spark.pyspark.virtualenv.enabled': 'true', 'spark.pyspark.virtualenv.type': 'native', 'spark.pyspark.virtualenv.bin.path': '/usr/bin/virtualenv'}, 'kind': 'pyspark'}
```

ID	YARN Application ID	Kind	State	Spark UI	Driver log	Current session?
2	application_1638027978840_0003	pyspark	idle	<a href="#">Link</a>	<a href="#">Link</a>	

In [2]:

```
sc.install_pypi_package("pandas==1.0.3")
sc.install_pypi_package("matplotlib==3.2.1")
sc.install_pypi_package("scipy==1.7.0")
sc.install_pypi_package("seaborn==0.10.0")
```

Starting Spark application

ID	YARN Application ID	Kind	State	Spark UI	Driver log	Current session?
3	application_1638027978840_0004	pyspark	idle			✓

SparkSession available as 'spark'.

```
Collecting pandas==1.0.3
  Using cached https://files.pythonhosted.org/packages/4a/6a/94b219b8ea0f2d580169e85ed1edc0163743f55aaeca8a44c2e8fc1e344
  e/pandas-1.0.3-cp37-cp37m-manylinux1_x86_64.whl
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.7/site-packages (from pandas==1.0.3)
Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib64/python3.7/site-packages (from pandas==1.0.3)
```

```
Collecting python-dateutil>=2.6.1 (from pandas==1.0.3)
  Using cached https://files.pythonhosted.org/packages/36/7a/87837f39d0296e723bb9b62bbb257d0355c7f6128853c78955f57342a56
d/python_dateutil-2.8.2-py2.py3-none-any.whl
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.6.1->pandas==1.0.3)
Installing collected packages: python-dateutil, pandas
Successfully installed pandas-1.0.3 python-dateutil-2.8.2

Collecting matplotlib==3.2.1
  Using cached https://files.pythonhosted.org/packages/b2/c2/71fcf957710f3ba1f09088b35776a799ba7dd95f7c2b195ec800933b276
b/matplotlib-3.2.1-cp37-cp37m-manylinux1_x86_64.whl
Requirement already satisfied: python-dateutil>=2.1 in /mnt/tmp/1638043908499-0/lib/python3.7/site-packages (from matplotlib==3.2.1)
Collecting pyparsing!=2.0.4,!>2.1.2,!>2.1.6,>=2.0.1 (from matplotlib==3.2.1)
  Using cached https://files.pythonhosted.org/packages/a0/34/895006117f6fce0b4de045c87e154ee4a20c68ec0a4c9a36d900888fb6b
c/pyparsing-3.0.6-py3-none-any.whl
Collecting cycler>=0.10 (from matplotlib==3.2.1)
  Using cached https://files.pythonhosted.org/packages/5c/f9/695d6bedebd747e5eb0fe8fad57b72fdf25411273a39791cde838d5a8f5
1/cycler-0.11.0-py3-none-any.whl
Requirement already satisfied: numpy>=1.11 in /usr/local/lib64/python3.7/site-packages (from matplotlib==3.2.1)
Collecting kiwisolver>=1.0.1 (from matplotlib==3.2.1)
  Using cached https://files.pythonhosted.org/packages/09/6b/6e567cb2e86d4e5939a9233f8734e26021b6a9c1bc4b1edccba236a84cc
2/kiwisolver-1.3.2-cp37-cp37m-manylinux_2_5_x86_64.manylinux1_x86_64.whl
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.1->matplotlib==3.2.1)
Installing collected packages: pyparsing, cycler, kiwisolver, matplotlib
Successfully installed cycler-0.11.0 kiwisolver-1.3.2 matplotlib-3.2.1 pyparsing-3.0.6

Collecting scipy==1.7.0
  Using cached https://files.pythonhosted.org/packages/b2/85/b00f13b52d079b5625e1a12330fc6453c947a482ff667a907c7bc60ed22
0/scipy-1.7.0-cp37-cp37m-manylinux_2_5_x86_64.manylinux1_x86_64.whl
Requirement already satisfied: numpy<1.23.0,>=1.16.5 in /usr/local/lib64/python3.7/site-packages (from scipy==1.7.0)
Installing collected packages: scipy
Successfully installed scipy-1.7.0

Collecting seaborn==0.10.0
  Using cached https://files.pythonhosted.org/packages/70/bd/5e6bf595fe6ee0f257ae49336dd180768c1ed3d7c7155b2fdf894c1c808
a/seaborn-0.10.0-py3-none-any.whl
Requirement already satisfied: pandas>=0.22.0 in /mnt/tmp/1638043908499-0/lib/python3.7/site-packages (from seaborn==0.10.0)
Requirement already satisfied: numpy>=1.13.3 in /usr/local/lib64/python3.7/site-packages (from seaborn==0.10.0)
Requirement already satisfied: scipy>=1.0.1 in /mnt/tmp/1638043908499-0/lib/python3.7/site-packages (from seaborn==0.10.0)
Requirement already satisfied: matplotlib>=2.1.2 in /mnt/tmp/1638043908499-0/lib/python3.7/site-packages (from seaborn==0.10.0)
```

```
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.7/site-packages (from pandas>=0.22.0->seaborn==0.10.0)
Requirement already satisfied: python-dateutil>=2.6.1 in /mnt/tmp/1638043908499-0/lib/python3.7/site-packages (from pandas>=0.22.0->seaborn==0.10.0)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /mnt/tmp/1638043908499-0/lib/python3.7/site-packages (from matplotlib>=2.1.2->seaborn==0.10.0)
Requirement already satisfied: cycler>=0.10 in /mnt/tmp/1638043908499-0/lib/python3.7/site-packages (from matplotlib>=2.1.2->seaborn==0.10.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /mnt/tmp/1638043908499-0/lib/python3.7/site-packages (from matplotlib>=2.1.2->seaborn==0.10.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/site-packages (from python-dateutil>=2.6.1->pandas>=0.22.0->seaborn==0.10.0)
Installing collected packages: seaborn
Successfully installed seaborn-0.10.0
```

```
In [3]: sc.list_packages()
```

Package	Version
beautifulsoup4	4.9.1
boto	2.49.0
click	7.1.2
cycler	0.11.0
jmespath	0.10.0
joblib	0.16.0
kiwisolver	1.3.2
lxml	4.5.2
matplotlib	3.2.1
mysqlclient	1.4.2
nltk	3.5
nose	1.3.4
numpy	1.16.5
pandas	1.0.3
pip	9.0.1
py-dateutil	2.2
pyparsing	3.0.6
python-dateutil	2.8.2
python37-sagemaker-pyspark	1.4.0
pytz	2020.1
PyYAML	5.3.1
regex	2020.7.14
scipy	1.7.0
seaborn	0.10.0

```
setuptools           28.8.0
six                  1.13.0
soup sieve           1.9.5
tqdm                4.48.2
wheel                0.29.0
windmill              1.6
```

## Importing

Now, import the installed packages from the previous block below.

```
In [4]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

## Loading Data

We are finally ready to load data. Using `spark` load the data from S3 into a `dataframe` object that we can manipulate further down in our analysis.

```
In [5]: df = spark.read.json('s3://sta9760-lei-yelpdatasets/yelp/yelp_academic_dataset_business.json')
```

```
In [6]: type(df)
```

```
<class 'pyspark.sql.dataframe.DataFrame'>
```

## Overview of Data

Display the number of rows and columns in our dataset.

```
In [7]: print(f'Total Columns: {len(df.dtypes)}')
print(f'Total Rows: {df.count():,}')
```

```
Total Columns: 14  
Total Rows: 160,585
```

In [8]:

```
df.printSchema()
```

```
root
|-- address: string (nullable = true)
|-- attributes: struct (nullable = true)
|   |-- AcceptsInsurance: string (nullable = true)
|   |-- AgesAllowed: string (nullable = true)
|   |-- Alcohol: string (nullable = true)
|   |-- Ambience: string (nullable = true)
|   |-- BYOB: string (nullable = true)
|   |-- BYOBCorkage: string (nullable = true)
|   |-- BestNights: string (nullable = true)
|   |-- BikeParking: string (nullable = true)
|   |-- BusinessAcceptsBitcoin: string (nullable = true)
|   |-- BusinessAcceptsCreditCards: string (nullable = true)
|   |-- BusinessParking: string (nullable = true)
|   |-- ByAppointmentOnly: string (nullable = true)
|   |-- Caters: string (nullable = true)
|   |-- CoatCheck: string (nullable = true)
|   |-- Corkage: string (nullable = true)
|   |-- DietaryRestrictions: string (nullable = true)
|   |-- DogsAllowed: string (nullable = true)
|   |-- DriveThru: string (nullable = true)
|   |-- GoodForDancing: string (nullable = true)
|   |-- GoodForKids: string (nullable = true)
|   |-- GoodForMeal: string (nullable = true)
|   |-- HairSpecializesIn: string (nullable = true)
|   |-- HappyHour: string (nullable = true)
|   |-- HasTV: string (nullable = true)
|   |-- Music: string (nullable = true)
|   |-- NoiseLevel: string (nullable = true)
|   |-- Open24Hours: string (nullable = true)
|   |-- OutdoorSeating: string (nullable = true)
|   |-- RestaurantsAttire: string (nullable = true)
|   |-- RestaurantsCounterService: string (nullable = true)
|   |-- RestaurantsDelivery: string (nullable = true)
|   |-- RestaurantsGoodForGroups: string (nullable = true)
|   |-- RestaurantsPriceRange2: string (nullable = true)
```

```
| -- RestaurantsReservations: string (nullable = true)
| -- RestaurantsTableService: string (nullable = true)
| -- RestaurantsTakeOut: string (nullable = true)
| -- Smoking: string (nullable = true)
| -- WheelchairAccessible: string (nullable = true)
| -- WiFi: string (nullable = true)
-- business_id: string (nullable = true)
-- categories: string (nullable = true)
-- city: string (nullable = true)
-- hours: struct (nullable = true)
| -- Friday: string (nullable = true)
| -- Monday: string (nullable = true)
| -- Saturday: string (nullable = true)
| -- Sunday: string (nullable = true)
| -- Thursday: string (nullable = true)
| -- Tuesday: string (nullable = true)
| -- Wednesday: string (nullable = true)
-- is_open: long (nullable = true)
-- latitude: double (nullable = true)
-- longitude: double (nullable = true)
-- name: string (nullable = true)
-- postal_code: string (nullable = true)
-- review_count: long (nullable = true)
-- stars: double (nullable = true)
-- state: string (nullable = true)
```

Display the DataFrame schema below.

Display the first 5 rows with the following columns:

- business\_id
- name
- city
- state
- stars
- categories

In [9]:

```
df.createOrReplaceTempView('business')
sqldf = spark.sql(
  ...
  SELECT business_id, name, city, state, stars, categories
  FROM business
  LIMIT 5
```

```
...  
)  
sqldf.show()
```

business_id	name	city	state	stars	categories
6iYb2HFDywm3zjuRg...	Oskar Blues Taproom	Boulder	CO	4.0	Gastropubs, Food,...
tCbdrrPZA0oiIYSmH...	Flying Elephants ...	Portland	OR	4.0	Salad, Soup, Sand...
bvN78f1M8NLprQ1a1...	The Reclamitory	Portland	OR	4.5	Antiques, Fashion...
oaepsyvc0J17qwi8c...	Great Clips	Orange City	FL	3.0	Beauty & Spas, Ha...
PE9uqAjdw0E4-8mjG...	Crossfit Terminus	Atlanta	GA	4.0	Gyms, Active Life...

## Analyzing Categories

Let's now answer this question: **how many unique categories are represented in this dataset?**

Essentially, we have the categories per business as a list - this is useful to quickly see what each business might be represented as but it is difficult to easily answer questions such as:

- How many businesses are categorized as `Active Life`, for instance
- What are the top 20 most popular categories available?

## Association Table

We need to "break out" these categories from the business ids? One common approach to take is to build an association table mapping a single business id multiple times to each distinct category.

For instance, given the following:

business_id	categories
abcd123	a,b,c

We would like to derive something like:

business_id	category
-------------	----------

business_id	category
abcd123	a
abcd123	b
abcd123	c

What this does is allow us to then perform a myriad of rollups and other analysis on this association table which can aid us in answering the questions asked above.

Implement the code necessary to derive the table described from your original yelp dataframe.

```
In [10]: df.select('business_id', 'categories').show(5)
```

```
+-----+-----+
| business_id | categories |
+-----+-----+
| 6iYb2HFDywm3zjuRg... | Gastropubs, Food, ... |
| tCbdrrPZA0oiIYSmH... | Salad, Soup, Sand... |
| bvN78f1M8NLprQ1a1... | Antiques, Fashion... |
| oaepsyvc0J17qwi8c... | Beauty & Spas, Ha... |
| PE9uqAjdw0E4-8mjG... | Gyms, Active Life... |
+-----+
only showing top 5 rows
```

```
In [11]: from pyspark.sql.functions import split, explode

association_table= df.select("business_id", explode(split(df.categories, ', '))\
                           .alias('category'))
association_table.show()
```

```
+-----+-----+
| business_id | category |
+-----+-----+
| 6iYb2HFDywm3zjuRg... | Gastropubs |
| 6iYb2HFDywm3zjuRg... | Food |
| 6iYb2HFDywm3zjuRg... | Beer Gardens |
| 6iYb2HFDywm3zjuRg... | Restaurants |
| 6iYb2HFDywm3zjuRg... | Bars |
```

```
| 6iYb2HFDywm3zjuRg...|American (Traditi...|
| 6iYb2HFDywm3zjuRg...|          Beer Bar|
| 6iYb2HFDywm3zjuRg...|        Nightlife|
| 6iYb2HFDywm3zjuRg...|       Breweries|
| tCbdrrPZA0oiIYSmH...|         Salad|
| tCbdrrPZA0oiIYSmH...|         Soup|
| tCbdrrPZA0oiIYSmH...| Sandwiches|
| tCbdrrPZA0oiIYSmH...|      Delis|
| tCbdrrPZA0oiIYSmH...| Restaurants|
| tCbdrrPZA0oiIYSmH...|      Cafes|
| tCbdrrPZA0oiIYSmH...| Vegetarian|
| bvN78f1M8NLprQ1a1...|   Antiques|
| bvN78f1M8NLprQ1a1...|      Fashion|
| bvN78f1M8NLprQ1a1...|      Used|
| bvN78f1M8NLprQ1a1...|Vintage & Consign...|
+-----+
only showing top 20 rows
```

Display the first 5 rows of your association table below.

```
In [12]: association_table.select('business_id','category').show(5)
```

```
+-----+
| business_id|    category|
+-----+
| 6iYb2HFDywm3zjuRg...| Gastropubs|
| 6iYb2HFDywm3zjuRg...|      Food|
| 6iYb2HFDywm3zjuRg...|Beer Gardens|
| 6iYb2HFDywm3zjuRg...| Restaurants|
| 6iYb2HFDywm3zjuRg...|      Bars|
+-----+
only showing top 5 rows
```

## Total Unique Categories

Finally, we are ready to answer the question: **what is the total number of unique categories available?**

Below, implement the code necessary to calculate this figure.

```
In [13]: association_table.select('category').distinct().count()
```

# Top Categories By Business

Now let's find the top categories in this dataset by rolling up categories.

## Counts of Businesses / Category

So now, let's unroll our distinct count a bit and display the per count value of businesses per category.

The expected output should be:

category	count
a	15
b	2
c	45

Or something to that effect.

```
In [14]: count_category= df.select(explode(split(df.categories, ', '))\
                                .alias('category'))\
.select('category').groupby('category').count().show()
```

category	count
Dermatologists	351
Paddleboarding	67
Aerial Tours	8
Hobby Shops	610
Bubble Tea	779
Embassy	9
Tanning	701
Handyman	507
Aerial Fitness	13
Falafel	141
Summer Camps	308

```

|      Outlet Stores| 184|
| Clothing Rental| 37|
| Sporting Goods| 1864|
| Cooking Schools| 114|
| College Counseling| 20|
| Lactation Services| 47|
| Ski & Snowboard S...| 55|
|          Museums| 336|
|           Doulas| 52|
+-----+
only showing top 20 rows

```

## Bar Chart of Top Categories

With this data available, let us now build a barchart of the top 20 categories.

**HINT:** don't forget about the matplotlib magic!

```
%matplotlib plt
```

```
In [15]: top_category = association_table.groupby('category')\
          .count()\
          .orderBy(['count'], ascending = False)
top_category.show(20)
```

```

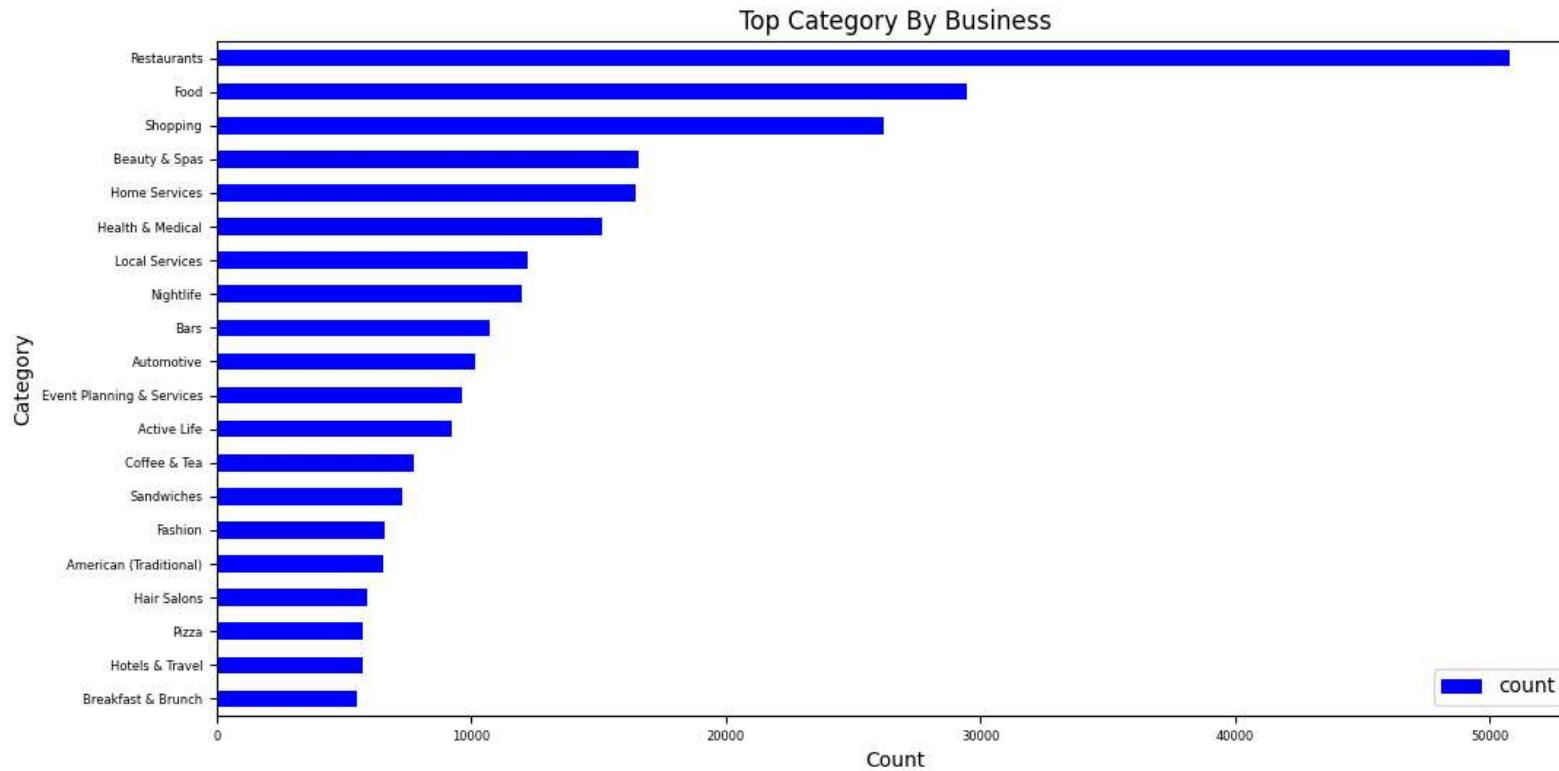
+-----+
|      category|count|
+-----+
| Restaurants|50763|
|      Food|29469|
|   Shopping|26205|
| Beauty & Spas|16574|
| Home Services|16465|
| Health & Medical|15102|
| Local Services|12192|
|   Nightlife|11990|
|       Bars|10741|
| Automotive|10119|
|Event Planning & ...| 9644|
| Active Life| 9231|
| Coffee & Tea| 7725|

```

```
|      Sandwiches| 7272|
|          Fashion| 6599|
|American (Traditi...| 6541|
|      Hair Salons| 5900|
|          Pizza| 5756|
|Hotels & Travel| 5703|
|Breakfast & Brunch| 5505|
+-----+-----+
only showing top 20 rows
```

```
In [16]: %matplotlib inline
```

```
In [17]: df_plt=top_category.toPandas().sort_values(by='count',ascending=False).head(20).sort_values(by='count',ascending=True) \
.plot(x='category',y='count',kind='barh',fontsize=6,figsize=(12, 6), color = 'blue')
plt.ylabel('Category')
plt.xlabel('Count')
plt.title('Top Category By Business')
%matplotlib plt
```



## Do Yelp Reviews Skew Negative?

Oftentimes, it is said that the only people who write a written review are those who are extremely *dissatisfied* or extremely *satisfied* with the service received.

How true is this really? Let's try and answer this question.

## Loading User Data

Begin by loading the user data set from S3 and printing schema to determine what data is available.

In [18]:

```
Reviewdf= spark.read.json('s3://sta9760-lei-yelpdatasets/yelp/yelp_academic_dataset_review.json')
```

```
Reviewdf.printSchema()
```

```
root
|-- business_id: string (nullable = true)
|-- cool: long (nullable = true)
|-- date: string (nullable = true)
|-- funny: long (nullable = true)
|-- review_id: string (nullable = true)
|-- stars: double (nullable = true)
|-- text: string (nullable = true)
|-- useful: long (nullable = true)
|-- user_id: string (nullable = true)
```

Let's begin by listing the `business_id` and `stars` columns together for the user reviews data.

```
In [19]: Reviewdf.createOrReplaceTempView('review')
sqldf = spark.sql(
    """
    SELECT business_id, stars
    FROM review
    LIMIT 5
    """
)
sqldf.show()
```

business_id	stars
DiRIdYhGyuTNZurKy...	3.0
bPmWDBkjBhV11Yk4B...	5.0
xHdKDNCJrvkYJkAGs...	4.0
Irp5sgl7XASH5ZTw2...	5.0
R8fLQ6TLz06MQR69K...	3.0

Now, let's aggregate along the `stars` column to get a resultant dataframe that displays *average stars* per business as accumulated by users who **took the time to submit a written review**.

```
In [20]: Reviewdf.createOrReplaceTempView('review')
sqldf = spark.sql(
    """
    SELECT business_id, AVG(stars) AS average_stars
    FROM review
    GROUP BY business_id
    ORDER BY average_stars DESC
    LIMIT 5
    """
)
sqldf.show()
```

```
SELECT business_id,avg(stars)
FROM review
GROUP BY business_id
LIMIT 5
...
)
sqldf.show()
```

business_id	avg(stars)
uEUweopM30lHcvxj0...	3.0
wdBrDCbZopowEkIEX...	4.538461538461538
L3WCfeVozu5etMhz4...	4.2
bOnsvrz1VkbrZM1jv...	3.8
R0IJhEI-zSJpYT1YN...	3.606060606060606

Now the fun part - let's join our two dataframes (reviews and business data) by `business_id`.

Let's see a few of these:

```
In [21]: df.createOrReplaceTempView('business')
Reviewdf.createOrReplaceTempView('review')
sqldf = spark.sql(
"""
SELECT DISTINCT AVG(review.stars) OVER (PARTITION BY review.business_id) AS avg_stars,business.stars,name,city,state
FROM
business
JOIN
review
ON business.business_id=review.business_id
ORDER BY avg_stars DESC
LIMIT 5
...
)
sqldf.show()
```

avg_stars	stars	name	city	state

5.0	5.0	Little Children S...	Brookline	MA
5.0	5.0	Gentle Touch Home...	Atlanta	GA
5.0	5.0	Dragonfly Preschool	Portland	OR
5.0	5.0	Reunion Car Care	Davenport	FL
5.0	5.0	Shamso Hair Studi...	Boston	MA

Compute a new dataframe that calculates what we will call the *skew* (for lack of a better word) between the avg stars accumulated from written reviews and the *actual* star rating of a business (ie: the average of stars given by reviewers who wrote an actual review **and** reviewers who just provided a star rating).

The formula you can use is something like:

```
(row['avg(stars)'] - row['stars']) / row['stars']
```

If the **skew** is negative, we can interpret that to be: reviewers who left a written response were more dissatisfied than normal. If **skew** is positive, we can interpret that to be: reviewers who left a written response were more satisfied than normal.

In [22]:

```
sqldf = spark.sql(
    """
    SELECT ROUND((avg_stars-stars)/stars,3) as skew
    FROM
    (SELECT DISTINCT AVG(review.stars) OVER (PARTITION BY review.business_id) AS avg_stars,business.stars,name,city,state
    FROM
    business
    JOIN
    review
    ON business.business_id=review.business_id
    )
    ...
)

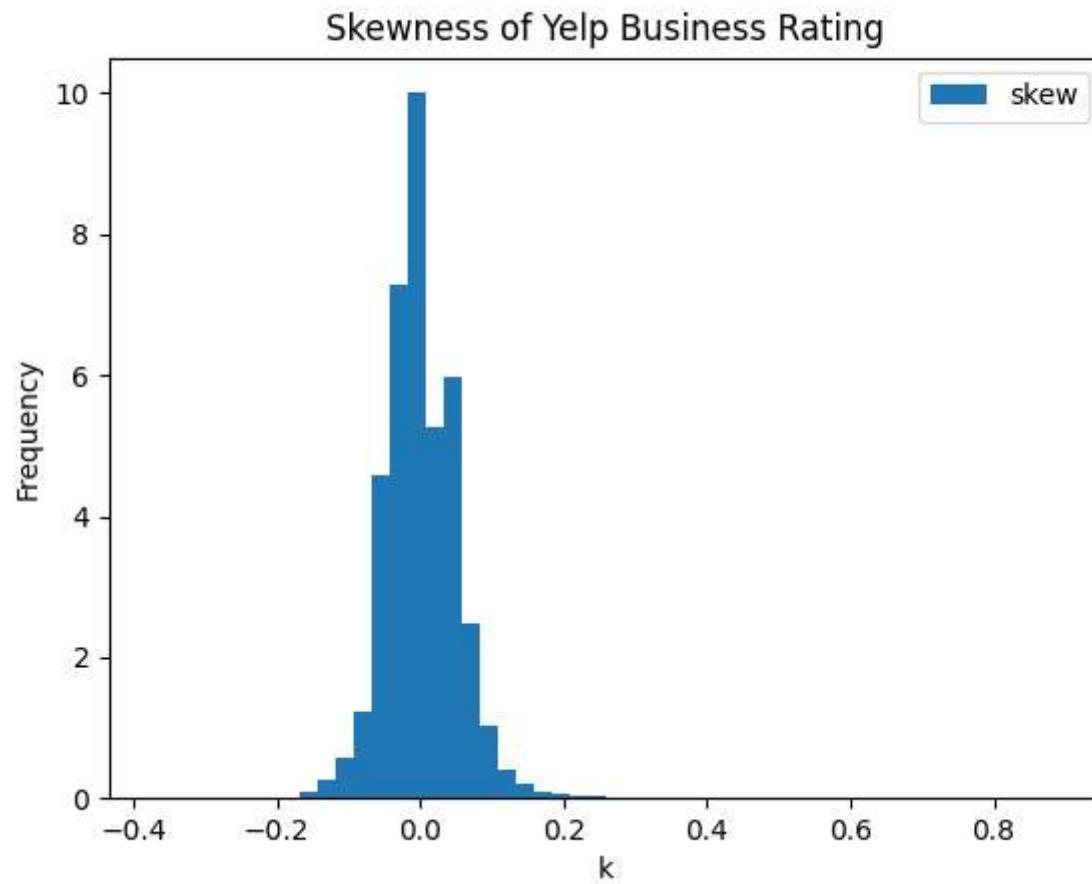
sqldf.show()
```

skew
0.0
0.0
0.123
0.026
0.007

```
| 0.014|  
|-0.043|  
| 0.0|  
| 0.014|  
| 0.086|  
|-0.022|  
| 0.086|  
|-0.044|  
|-0.042|  
| 0.056|  
|-0.111|  
| 0.067|  
|-0.049|  
| 0.083|  
| 0.05|  
+-----+  
only showing top 20 rows
```

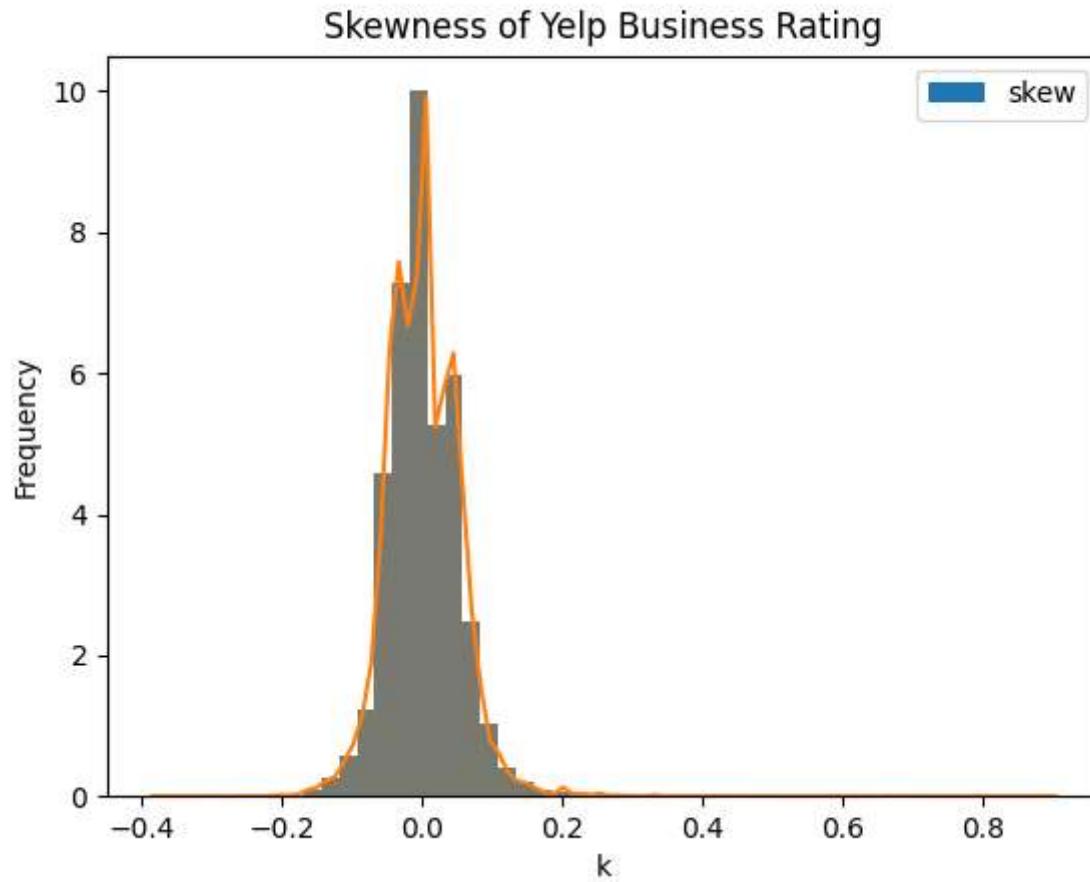
In [23]:

```
sqldf.toPandas().plot(kind='hist',bins=50,title = "Skewness of Yelp Business Rating",density=True)  
plt.xlabel('k')  
%matplotlib plt
```



And finally, graph it!

```
In [24]: sns.distplot(sqldf.toPandas(), bins =50).  
        set(title = "Skewness of Yelp Business Rating")  
  
%matplotlib plt
```



So, do Yelp (written) Reviews skew negative? Does this analysis actually prove anything? Expound on implications / interpretations of this graph.

Yes, it is skew negative. Reviewers who left a written response were more unsatisfied than normal

## Should the Elite be Trusted? (Or, some other analysis of your choice)

For the final portion - you have a choice:

- Try and analyze some interesting dimension to this data. The **ONLY** requirement is that you must use the **Users** dataset and join on either the **business\*** or **reviews\*\*** dataset
- Or, you may try and answer the question posed: how accurate or close are the ratings of an "elite" user (check Users table schema) vs the actual business rating.

Feel free to use any and all methodologies at your disposal - only requirement is you must render one visualization in your analysis

## Loading user data

Begin by loading the user data set from S3 and printing schema to determine what data is available.

```
In [25]: Userdf=spark.read.json('s3://sta9760-lei-yelpdatasets/yelp/yelp_academic_dataset_user.json')
```

```
In [26]: Userdf.printSchema()
Reviewdf.printSchema()
```

```
root
|-- average_stars: double (nullable = true)
|-- compliment_cool: long (nullable = true)
|-- compliment_cute: long (nullable = true)
|-- compliment_funny: long (nullable = true)
|-- compliment_hot: long (nullable = true)
|-- compliment_list: long (nullable = true)
|-- compliment_more: long (nullable = true)
|-- compliment_note: long (nullable = true)
|-- compliment_photos: long (nullable = true)
|-- compliment_plain: long (nullable = true)
|-- compliment_profile: long (nullable = true)
|-- compliment_writer: long (nullable = true)
|-- cool: long (nullable = true)
|-- elite: string (nullable = true)
|-- fans: long (nullable = true)
|-- friends: string (nullable = true)
|-- funny: long (nullable = true)
|-- name: string (nullable = true)
|-- review_count: long (nullable = true)
|-- useful: long (nullable = true)
```

```
|-- user_id: string (nullable = true)
|-- yelping_since: string (nullable = true)

root
|-- business_id: string (nullable = true)
|-- cool: long (nullable = true)
|-- date: string (nullable = true)
|-- funny: long (nullable = true)
|-- review_id: string (nullable = true)
|-- stars: double (nullable = true)
|-- text: string (nullable = true)
|-- useful: long (nullable = true)
|-- user_id: string (nullable = true)
```

In [27]:

```
Userdf.createOrReplaceTempView('user')
Reviewdf.createOrReplaceTempView('review')
user_t = spark.sql(
    """
        SELECT user_id, elite, average_stars, review_count
        FROM user
    """
)
review_t = spark.sql(
    """
        SELECT user_id,business_id,stars
        FROM review
    """
)
user_t.show()
review_t.show()
```

user_id	elite	average_stars	review_count
q_QQ5kBBlCcbL1s4...	2006,2007,2008,20...	3.85	1220
dIIKEfOgo0KqUfGQv...	2007,2008,2009,20...	4.09	2136
D6ErcUnFALnCQN4b1...	2010,2011	3.76	119
JnPIjvC0cmooNDFsa...	2009,2010,2011,20...	3.77	987
37Hc8hr3cw0iHLoPz...	2009,2010,2011	3.72	495
n-QwITZYrXlKQRiV3...	2010,2011,2012,20...	3.59	229
eCJoZqpV1fDKJGAsX...		3.86	51
cojec0wQJpsYDxnjt...	2010,2011	3.79	51
1jXmzuIFKxTnEnR0p...	2010,2011,2012,20...	3.43	299

-8Qo0IfvwvxJ4sY20...	2012, 2013, 2014, 20...	3.88	288
EtofuImujQBS002xa...	2009, 2010	3.83	44
cxS6dbjyPgPS1S890...		4.18	65
MUzkXfPS9JaMgJ907...	2010, 2012, 2013, 20...	4.16	750
tjwb1GkWN9m0vsGay...	2010, 2011, 2012, 20...	3.66	632
m-zIVssiXN4bnDFqM...		4.5	2
fxqvyXlml4400Bgl...	2011, 2012, 2013, 20...	3.65	363
9edAbpnihHFdpAvk...	2007, 2008, 2009, 2010	3.73	356
wURnB9fRNNGAl13yB...	2009, 2010, 2011, 2012	3.75	77
14P65LXNBnJqI7oTX...	2008, 2009, 2010, 20...	3.93	202
9RIX1hUb_xEVuc_o0...		3.57	53

only showing top 20 rows

user_id	business_id	stars
ak0TdVmGKo4pwqdJS...	buF9druCkbuXLX526...	4.0
YoVfDbnIS1W0f7abN...	RA4V8pr014UyUbDvI...	4.0
eC5evKn1TWDyHCyQA...	_sS2LBIGNT5NQb6PD...	5.0
SFQ1jcнGgu00LYWnb...	0AzLzHfoJgL7R0whd...	2.0
0kA0PAJ8QFMeveQWH...	8zehGz9jnxPqXt0c7...	4.0
RNm_RWkcd02Li2mKP...	xGXzsc-hzam-VArK6...	1.0
Q8c91v7luItVB0cMF...	EXOsmAB1s71WePlQk...	2.0
XGkAG92TQ3MQUKGX9...	DbXHN1890xSXNiYRc...	5.0
LWUnzwK0ILquLLZcH...	mD-A9KOWADXvfrZfw...	4.0
99RsBrARhhx60UnAC...	EEHhKSxUvJkoPSzeG...	5.0
eLAYHxHUutiXswy-C...	WQFn1A7-UAA4JT5YW...	1.0
Ng183gs3n22SzLASN...	iZxfwlfnPXPV_0S1x0...	2.0
hn0ZbitvmlHnF--KJ...	TA1KUSCu8GkWP9w0r...	4.0
B7YSV6r1ePAXc69Fk...	wZgUAuDuEGPEzKK-P...	1.0
xpxWG7jQXZE6BcSeu...	mw_qxZJraNu7Q6u0G...	5.0
HvpNr0ohHCaVLp014...	gGvNgShksetPoimyK...	4.0
bUHweiErUJ36WGeNr...	6Hm2FmfLcU_M91TrZ...	5.0
JHXQEayrDH0WGexs0...	KXCXaF5qimmtKKqnP...	1.0
DECuRZwkUw8ELQzfN...	zmZ3HkVCeZPBefJJx...	4.0
jySmPCKEkJR3cWJlk...	ZW7aI5F0_3q_vSzI4...	5.0

only showing top 20 rows

## Join user table with reviews table

In [28]: `user_t.createOrReplaceTempView('user_t')`

```

review_t.createOrReplaceTempView('review_t')
UserReviewdf = spark.sql(
    """
    SELECT review_t.business_id,review_t.stars,user_t.elite,user_t.average_stars,user_t.review_count
    FROM review_t,user_t
    WHERE review_t.user_id=user_t.user_id
    """
)
UserReviewdf.show()

```

business_id	stars	elite	average_stars	review_count
GgR7kcKykuqXB11fW...	5.0		2.62	12
rxNfidGLHtMMyLNeo...	5.0		3.67	11
bAuY0a-VuqTOnKzWN...	2.0		2.73	11
vqQXI-Pxz3izeTUF6...	5.0		2.73	11
20aX6XjAoI7VD6jLd...	2.0		2.73	11
Ifoj3AxP13Exsd_Yl...	1.0		2.73	11
piuF4DdprICZgRent...	2.0		3.67	11
_ihrl49Ua40_qd_G-...	5.0		3.67	11
huvei0ZFjsFI_Aiq9...	4.0		3.67	11
ggqr9BEklylICK177...	2.0		3.67	11
kj77ZrCsFlVuqWkIr...	5.0		3.67	11
szJPyyYqJRxithtH6...	4.0		3.67	11
VZVkkZx2Zsq0ES8Sfo...	1.0		2.25	7
ENBeXC1MunFTlgl_7...	1.0		2.25	7
Tav-90WqGKFestUqt...	4.0		4.75	2
OMOxwfTxPilOunHS2...	5.0		4.43	7
KAwB6KyjhUizBhFo8...	5.0		4.43	7
2kKbqNPY6j1AbpXxY...	5.0		4.43	7
qUD80opa9pUTPVudH...	5.0		4.43	7
f-dvot7GwmnObBdCl...	5.0		4.43	7

only showing top 20 rows

## Replace Elite user with YES, otherwise NO

In [29]:

```

UserReviewdf.createOrReplaceTempView('ur')
elite_user = spark.sql(
    """
    """
)

```

```

SELECT business_id,stars,average_stars, review_count, (CASE
WHEN elite like '' THEN 'NO' ELSE 'YES' END)AS Elite_Y_N
FROM ur
...
)

elite_user.show()

```

business_id	stars	average_stars	review_count	Elite_Y_N
GgR7kcKykqXB11fw...	5.0	2.62	12	NO
rxNfidGLHtMYyLNeo...	5.0	3.67	11	NO
bAuY0a-VuqTOnKzWN...	2.0	2.73	11	NO
vqQXI-Pxz3izeTUF6...	5.0	2.73	11	NO
20aX6XjAoI7VD6jLd...	2.0	2.73	11	NO
If0j3AxP13Exsd_Y1...	1.0	2.73	11	NO
szJPyyYqJRxithH6...	4.0	3.67	11	NO
piuF4DdprICZgRent...	2.0	3.67	11	NO
_ihrL49Ua40_qd_G-...	5.0	3.67	11	NO
ggqr9BEklylICK177...	2.0	3.67	11	NO
kj77ZrCsFlVuqWkIr...	5.0	3.67	11	NO
huveiOZFjsFI_Aiq9...	4.0	3.67	11	NO
VZVkJx2Zsq0ES8Sfo...	1.0	2.25	7	NO
ENBeXC1MunFT1gl_7...	1.0	2.25	7	NO
Tav-90WqGKFestUqt...	4.0	4.75	2	NO
OMOxwfTxPil0unHS2...	5.0	4.43	7	NO
wpm_IUt_AoZOXdUvP...	5.0	4.43	7	NO
2kKbqNPY6j1AbpXxY...	5.0	4.43	7	NO
qUD80opa9pUTPVudH...	5.0	4.43	7	NO
f-dvot7GwmnObBdCl...	5.0	4.43	7	NO

only showing top 20 rows

## Calculate the difference between all users' ratings with actual business ratings

In [30]:

```

elite_user.createOrReplaceTempView('elite_user')
star_diff = spark.sql(
...
SELECT business_id,stars,average_stars,abs(stars-average_stars) AS diff, Elite_Y_N

```

```
FROM elite_user
...
)
star_diff.show()
```

business_id	stars	average_stars	diff	Elite_Y_N
GgR7kcKykuqXB11fW...	5.0	2.62	2.38	NO
rxNfidGLHtMYyLNeo...	5.0	3.67	1.33	NO
bAuY0a-VuqTOnKzWN...	2.0	2.73	0.73	NO
vqQXI-Pxz3izeTUF6...	5.0	2.73	2.27	NO
20aX6XjAoI7VD6jLd...	2.0	2.73	0.73	NO
If0j3AxPl3Exsd_Yl...	1.0	2.73	1.73	NO
piuF4DdprICZgRent...	2.0	3.67	1.67	NO
_ihrl49Ua40_qd_G-...	5.0	3.67	1.33	NO
huveiOZFjsFI_Aiq9...	4.0	3.67	0.33000000000000007	NO
szJPyyYqJRxithH6...	4.0	3.67	0.33000000000000007	NO
ggqr9BEklylICK177...	2.0	3.67	1.67	NO
kj77ZrCsFlVuqWkIr...	5.0	3.67	1.33	NO
VZVkJx2Zsq0ES8Sfo...	1.0	2.25	1.25	NO
ENBeXC1MunFT1gl_7...	1.0	2.25	1.25	NO
Tav-90WqGKFestUqt...	4.0	4.75	0.75	NO
2kKbqNPY6j1AbpXxY...	5.0	4.43	0.5700000000000003	NO
f-dvot7GwmnObBdC1...	5.0	4.43	0.5700000000000003	NO
OMOxwfTxPilOunHS2...	5.0	4.43	0.5700000000000003	NO
KAwB6KyjhUizBhFo8...	5.0	4.43	0.5700000000000003	NO
wpm_IUt_AoZOXdUvP...	5.0	4.43	0.5700000000000003	NO

only showing top 20 rows

**Extract and filter users are elit, display their ratings difference with actual business ratings**

```
In [31]: star_diff.createOrReplaceTempView('star_diff')
elit_diff = spark.sql(
...
SELECT diff, Elite_Y_N
FROM star_diff
WHERE Elite_Y_N="YES"
...)
```

```
)  
elit_diff.show()
```

```
+-----+-----+  
|       diff|Elite_Y_N|  
+-----+-----+  
| 0.16999999999999993|     YES|  
| 0.5299999999999998|     YES|  
|      1.35|     YES|  
|0.00999999999999787|     YES|  
| 0.370000000000001|     YES|  
| 0.370000000000001|     YES|  
| 0.990000000000002|     YES|  
| 0.850000000000001|     YES|  
| 0.370000000000001|     YES|  
| 1.819999999999998|     YES|  
| 0.149999999999999|     YES|  
| 0.589999999999999|     YES|  
| 1.180000000000002|     YES|  
| 0.629999999999999|     YES|  
| 0.870000000000001|     YES|  
| 0.990000000000002|     YES|  
| 0.100000000000009|     YES|  
|      1.37|     YES|  
| 0.370000000000001|     YES|  
|      1.35|     YES|  
+-----+-----+  
only showing top 20 rows
```

## Visualization of elit users rating with actual business star ratings

Now, Following is the histogram displays absolute difference between the ratings of elit users and the actual business star ratings. The difference is centralizing between [-1.25-1.25].There is slight rating difference. But overall, the rating of the elite should be trusted.

```
In [32]: elit_diff.createOrReplaceTempView('elit_diff')  
diff = spark.sql(  
    ...  
    "SELECT diff  
    FROM elit_diff  
    ...  
)
```

```
diff.toPandas().plot(kind='hist',bins=50,density=True)
plt.xlabel('abs_diff')
%matplotlib plt
```

